

# Advanced Programming II

## Lesson 3: Concurrent Programming viewed as an Abstraction

Grado en Ingeniería Informática

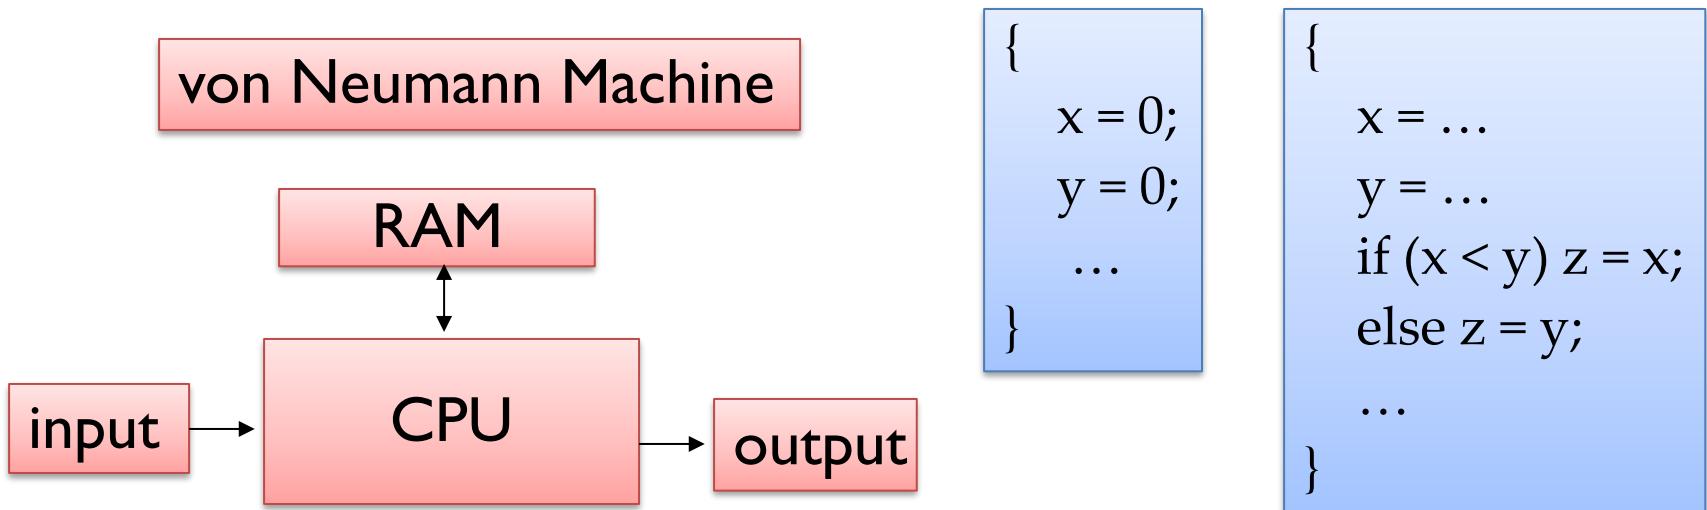
Grado en Ingeniería del Software

# Contents

- From Sequential programming to Concurrent programming
- Advantages and usages of Concurrent programming
- Inherent problems in Concurrent programming:
  - Atomic sentences
  - Critical Sections
  - Mutual Exclusion
- Platforms for concurrent execution

# Sequential Programming

- To take the most of the computer's power by means of Sequential Programming languages, a well-suited program must be explicitly created
- For each clock cycle:
  - An instruction is fetched from memory
  - It is decodified and executed by sending several electronic signals to the corresponding system components



# Formal Example

- If  $P = p_1; p_2; \dots; p_n$  is a sequential program
  - $p_1$  executes always before  $p_2$  (in any execution),
  - $p_2$  executes always before  $p_3$
  - ...
  - $p_{n-1}$  executes always before  $p_n$
- Stated that the symbol " $\rightarrow$ " means "precedes to" and  $\forall e$  means "for any execution", the sequential behaviour may be formalized as  $\forall e. (p_1 \rightarrow p_2) \wedge (p_2 \rightarrow p_3) \wedge \dots \wedge (p_{n-1} \rightarrow p_n)$
- This behaviour is always the same, even if the code  $P$  includes conditional or loop sentences.

**if (b) {A;} else {B;}**

- $\forall e. (b \rightarrow A) \vee (b \rightarrow B)$  (depends on input data for each execution)

# Formal Example

**while (b) {A;}**

$\forall e.b \vee (b \rightarrow A \rightarrow b) \vee (b \rightarrow A \rightarrow b \rightarrow A \rightarrow b) \vee \dots$

(The number of times **b** is executed depends on input)

Let's say  $P = p_1; p_2; \dots; p_n$

assuming  $0 < i < j \leq n$  then  $\forall e.p_i \rightarrow p_j$

The program is **deterministic**. Given a program and an input data set, it can be predicted which is the next sentence to execute

Sentences are strictly ordered

# Concurrent Programming

- Not all the sentences of a program **have to** be executed sequentially.

```
{  
    x = 0;  
    y = 0;  
    z = 0;  
    ...  
}
```

The language requires the programmer to set an order of execution

# Concurrent Programming

- Not all the sentences of a program have to be executed sequentially
  - This code can be executed following **6** different orders, and any of them is considered as valid

```
{  
    x = 0;  
    y = 0;  
    z = 0;  
    ...  
}
```

```
{  
    x = 0;  
    z = 0;  
    y = 0;  
    ...  
}
```

```
{  
    y = 0;  
    x = 0;  
    z = 0;  
    ...  
}
```

.....

# Concurrent Programming

- Not all the sentences of a program have to be executed sequentially
- This code can be executed following 6 different orders, and any of them is considered valid
- Assume we define the **new operator ||** to express formally this behaviour

```
x = 0 || y = 0 || z = 0    = {x = 0 → y = 0 → z = 0,  
                                x = 0 → z = 0 → y = 0,  
                                y = 0 → x = 0 → z = 0,  
                                ...  
}
```

# Concurrent Programming

- Assume we define the new operator  $\parallel$  to express formally this behaviour.

But, what happens if we have 3 processors?

- Each block of code (each sentence) can be assigned to a different processor, achieving a possibly higher throughput and a valid execution.
- The meaning of **P  $\parallel$  Q is extended** to express that there is a valid execution of P and Q where both blocks of code overlap in time:

$$\exists \text{ e. } \neg(P \rightarrow Q) \wedge \neg(Q \rightarrow P)$$

# Concurrent Programming

- How to understand well the meaning of an execution like  $P \parallel Q$ ?
  - We may require two real processors in order to have a true time overlapping processing of the code
    - It is not a good idea because, in such case, the behaviour of the code depends on the hardware architecture
  - Instead of this, we will assume that there are a couple of **virtual processors**, and each of them executes a block of code
    - In the real world, we could assign a real processor to each virtual one, but what happens if there are not enough real processors?

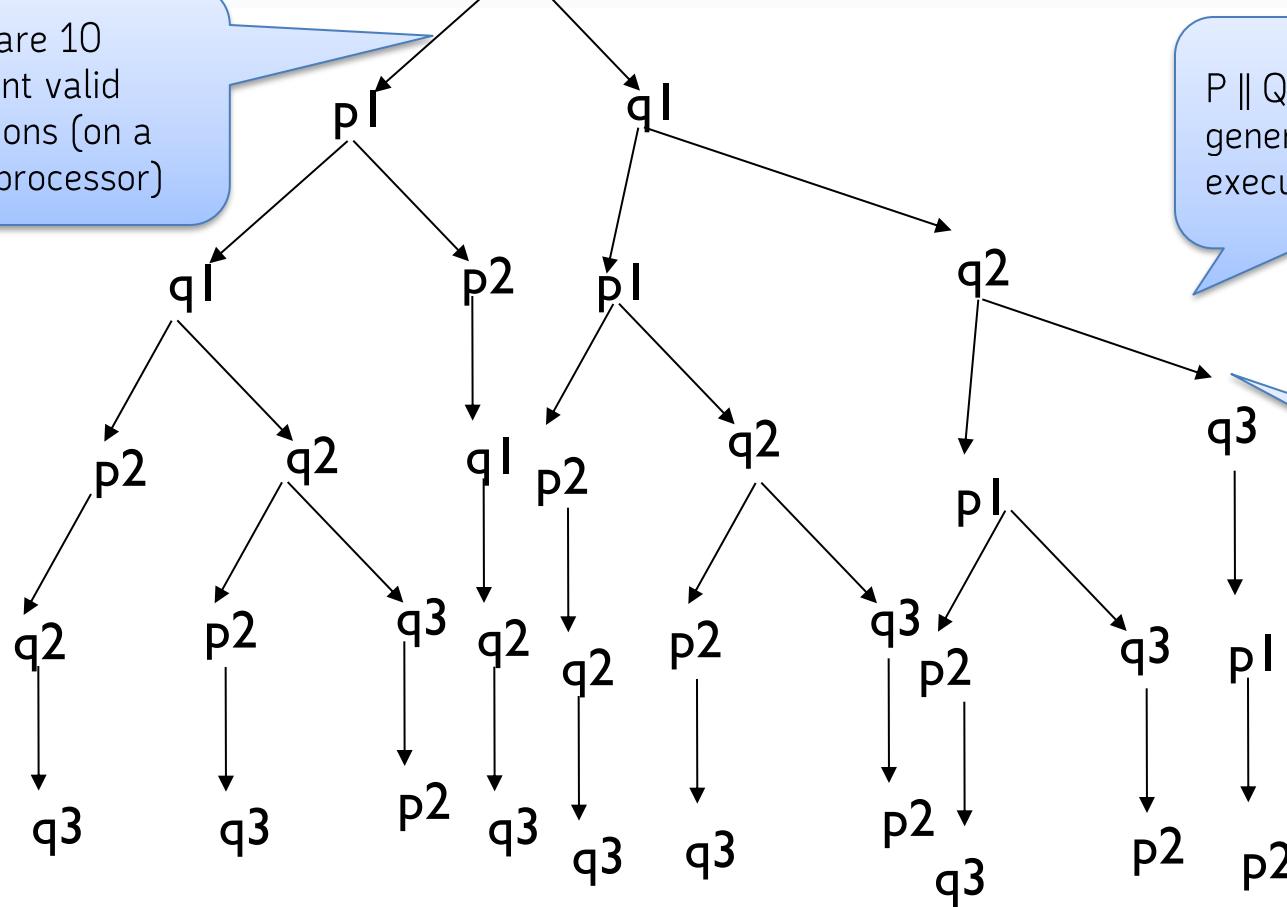
Let's say  $P = p_1 \rightarrow p_2 \rightarrow \dots \rightarrow p_n$  and  
 $Q = q_1 \rightarrow q_2 \rightarrow \dots \rightarrow q_m$

How can be interpreted  $P \parallel Q$  when there is only a single real processor?

# Interleaving

- Let's say  $n = 2$  and  $m = 3$ :  $P = p_1 \rightarrow p_2$  and  $Q = q_1 \rightarrow q_2 \rightarrow q_3$

There are 10 different valid executions (on a single processor)



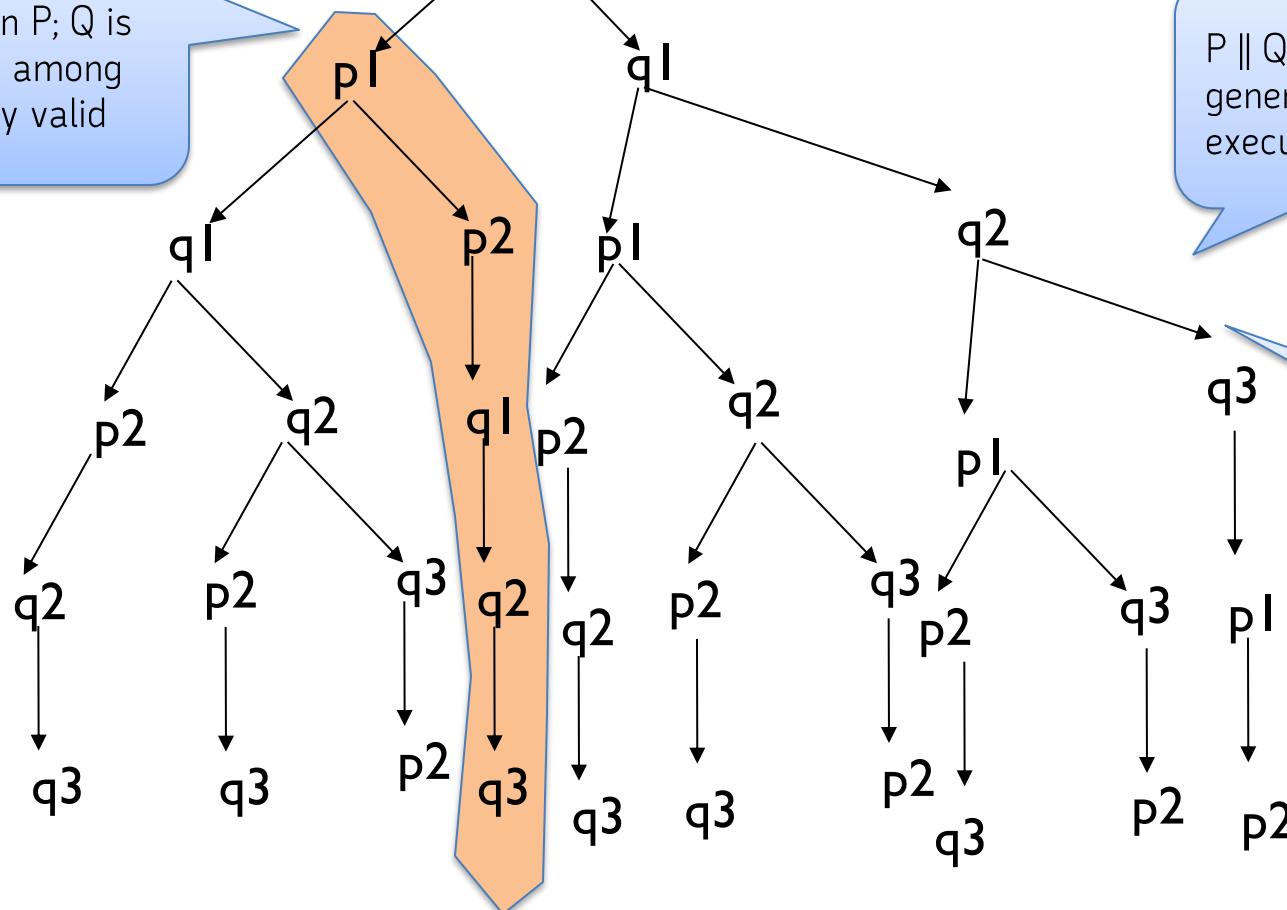
$P \parallel Q$   
generates an  
execution tree

Each branch of  
the tree  
constitutes a valid  
execution of  $P \parallel Q$

# Interleaving

- Let's say  $n = 2$  and  $m = 3$ :  $P = p_1 \rightarrow p_2$  and  $Q = q_1 \rightarrow q_2 \rightarrow q_3$

The sequential execution  $P; Q$  is only one among the many valid ones



$P \parallel Q$  generates an execution tree

Each branch of the tree constitutes a valid execution of  $P \parallel Q$

# Partial Order

- A valid execution of  $P \parallel Q$   
 $(P = p_1 \rightarrow p_2 \rightarrow \dots \rightarrow p_n, Q = q_1 \rightarrow q_2 \rightarrow \dots \rightarrow q_m)$   
is obtained by interspersing P and Q sentences,  
**but not all merges are valid**
- The internal order of the sentences in P and Q must remain,  
Given two indexes i, j,
  - if  $i < j$  then  $\forall e. (p_i \rightarrow p_j) \wedge (q_i \rightarrow q_j)$
  - if  $i < j$  then  $\forall e. (p_i \rightarrow q_j) \vee (q_j \rightarrow p_i)$

The order inside P  
and Q is kept up

But there is no need of  
any order among the  
sentences of P and Q

# Partial Order

- A valid execution of  $P \parallel Q$   
 $(P = p_1 \rightarrow p_2 \rightarrow \dots \rightarrow p_n, Q = q_1 \rightarrow q_2 \rightarrow \dots \rightarrow q_m)$   
is obtained by interspersing P and Q sentences,  
**but not all merges are valid**
- The internal order of the sentences in P and Q must remain,  
Given two indexes i, j,
  - if  $i < j$  then  $\forall e. (p_i \rightarrow p_j) \wedge (q_i \rightarrow q_j)$
  - if  $i < j$  then  $\forall e. (p_i \rightarrow q_j) \vee (q_j \rightarrow p_i)$

At any time, different sentences are available to be executed and one of them is randomly selected. Therefore, a parallel execution is always **NON DETERMINISTIC**

The sentences of  $P \parallel Q$  are **partially ordered**

# Concurrent Programming

- How to understand well the meaning of an execution like **P || Q**?
  - We may require two real processors in order to have a true time overlapping processing of the code
    - It is not a good idea because, in such a case, the behaviour of the code depends on the hardware architecture
  - Instead of this, we will assume that here are a couple of **virtual processors**, and each of them executes a block of code
  - A virtual processor may correspond with a real one, but
    - what happens if there are not enough real processors?
  - The sentences of P and Q are interspersed in a processor
    - (semantics of interleaving)
- This may be seen as if the virtual processors executing P and Q are working with different speeds, so
  - We can not take for granted anything about the relative speed of virtual processors, or about the exact time of execution taken by any block of code

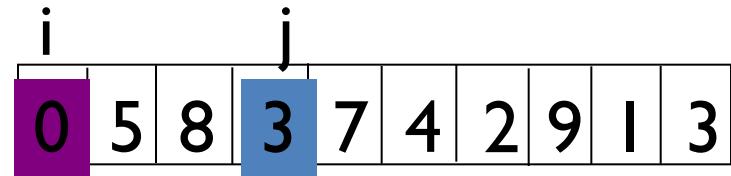
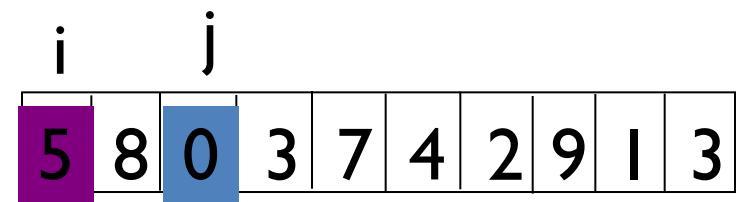
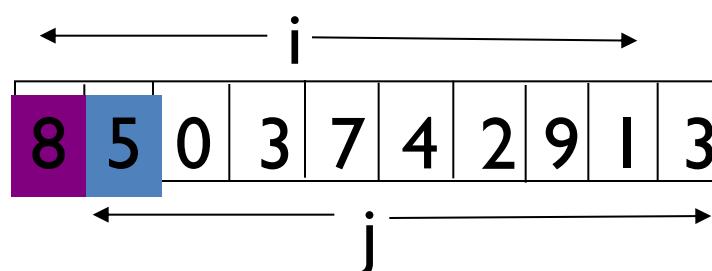
# Concurrent Programming

Summary,

- $P \parallel Q$  stands for a concurrent execution of P and Q
- P and Q are named **processes/threads** because they are executed concurrently with other processes
- Nonetheless, the code of P and Q are executed **sequentially**
- If  $P = p_1 \rightarrow p_2 \rightarrow \dots \rightarrow p_n$  and  $Q = q_1 \rightarrow q_2 \rightarrow \dots \rightarrow q_m$ , then the total number of valid executions of **P||Q** (assuming no overlapping in time when two sentences are executed) is
$$(m+n)!/m!*n!$$
- Each valid execution is named **trace** (a branch of the tree)
- A program **P||Q** is **correct** when all its traces are correct

# Advantages

- Improve the throughput of processors
- Exploit multiprocessor architectures
- Make easier to model systems where the coherence is inherent
- Complete a task quickly
  - Example: Selection Sort



# Example (Selection Sort)

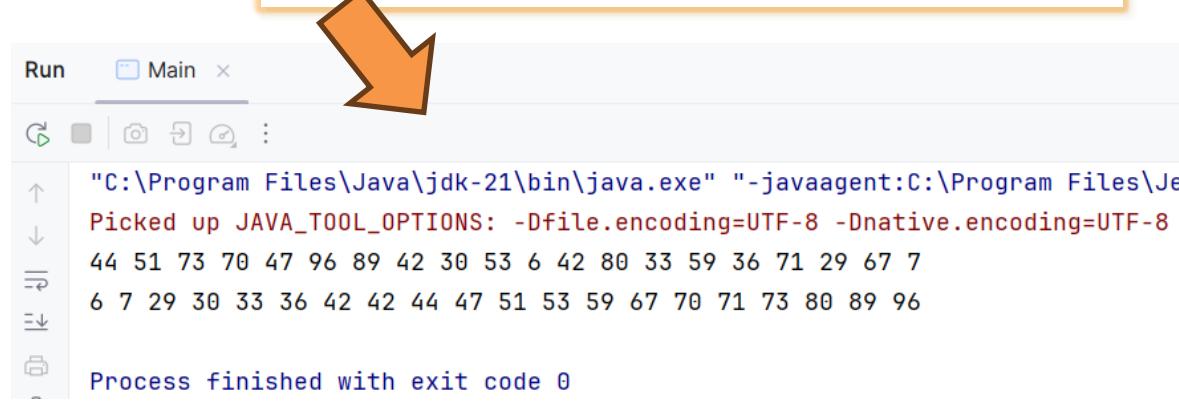
Implementing the Selection sort algorithm

```
def selectionSort(lista: Array[Int]) =  
    for (i<-0 until lista.length)  
        for (j<-i+1 until lista.length)  
            if (lista(i) > lista(j))  
                val aux = lista(i)  
                lista(i) = lista(j)  
                lista(j) = aux
```

# Example (Selection Sort). Option 1

Sequential case: ordering an array by calling the previous algorithm

```
import scala.util.Random
object Main extends App:
    val a = new Array[Int](20)
    for (i <- 0 until a.length)
        a(i) = Random.nextInt(100)
    println(a.mkString(" "))
    selectionSort(a)
    println(a.mkString(" "))
```



The screenshot shows a Java IDE interface with a run window titled "Run". Inside the run window, there is a tab labeled "Main" which is currently selected. Below the tabs, there are several icons for file operations like opening, saving, and running. The main content area displays the output of the Java application. The output starts with the command used to run the application: "C:\Program Files\Java\jdk-21\bin\java.exe" "-javaagent:C:\Program Files\Java\ToolBox\agent.jar". It then shows the message "Picked up JAVA\_TOOL\_OPTIONS: -Dfile.encoding=UTF-8 -Dnative.encoding=UTF-8". Following this, the initial state of the array is printed as a string of numbers separated by spaces: "44 51 73 70 47 96 89 42 30 53 6 42 80 33 59 36 71 29 67 7". Below this, the sorted state of the array is printed: "6 7 29 30 33 36 42 42 44 47 51 53 59 67 70 71 73 80 89 96". At the bottom of the output, the message "Process finished with exit code 0" is displayed.

# Example (Selection Sort). Option 2

Use the divide and conquer technique: the array is divided into two halves. Then, we call the algorithm twice to sort each half. Finally, we need a method that merges the two halves already ordered into the definitive sorted array.

```
def mergeIt(l1: Array[Int], l2: Array[Int], lista: Array[Int]) =  
    var i=0  
    var j=0  
    var k=0  
    while i<l1.length && j<l2.length do  
        if (l1(i)<=l2(j)) then  
            lista(k) = l1(i); i += 1  
        else  
            lista(k) = l2(j); j += 1  
        k += 1  
    while i < l1.length do  
        lista(k) = l1(i); i += 1; k += 1  
    while j < l2.length do  
        lista(k) = l2(j); j += 1; k += 1
```

# Example (Selection Sort). Option 2

Use the divide and conquer technique: the array is divided into two halves. Then, we call the algorithm twice to sort each half. Finally, we need a method that merges the two halves already ordered into the definitive sorted array.

```
import scala.util.Random
object Main extends App:
    val a = new Array[Int](20)
    for (i <- 0 until a.length)
        a(i) = Random.nextInt(100)
    println(a.mkString(" "))
    val (l1, l2) = a.splitAt(a.length / 2)

    selectionSort(l1)
    selectionSort(l2)
    mergelt(l1, l2, a)

    println(a.mkString(" "))
```



```
"C:\Program Files\Java\jdk-21\bin\java.exe" "-javaagent:C:\I
Picked up JAVA_TOOL_OPTIONS: -Dfile.encoding=UTF-8 -Dnative
48 68 66 81 9 95 55 39 63 87 95 49 13 82 62 24 92 68 15 38
9 13 15 24 38 39 48 49 55 62 63 66 68 68 81 82 87 92 95 95
Process finished with exit code 0
```

# Example (Selection Sort). Option 3

With **parallelism**: using **threads** + divide and conquer. Two threads are used **simultaneously** to sort the halves **at the same time**. Finally, both halves are merged.

```
import scala.util.Random
object Main extends App:
    val a = new Array[Int](20)
    for (i <- 0 until a.length)
        a(i) = Random.nextInt(100)
    println(a.mkString(" "))
    val (a1, a2) = a.splitAt(a.length / 2)
    val h1 = new MyThread(a1)
    val h2 = new MyThread(a2)
    h1.start();
    h2.start()
    h1.join();
    h2.join()

    mergelt(a1, a2, a)
    println(a.mkString(" "))
```



```
"C:\Program Files\Java\jdk-21\bin\java.exe" "-javaagent:C:\Picked up JAVA_TOOL_OPTIONS: -Dfile.encoding=UTF-8 -Dnative
42 58 81 77 44 19 73 94 87 35 73 25 71 59 72 12 10 8 16 61
8 10 12 16 19 25 35 42 44 58 59 61 71 72 73 73 77 81 87 94
```

```
Process finished with exit code 0
```

# Example

- Studying complexities

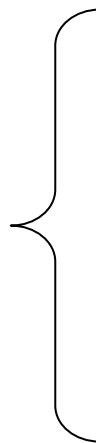
```
def selectionSort(lista: Array[Int]) =  
    for (i<-0 until lista.length)  
        for (j<-i+1 until lista.length)  
            if (lista(i) > lista(j))  
                val aux = lista(i)  
                lista(i) = lista(j)  
                lista(j) = aux
```

If  $n$  is the number of items in the vector to sort

$$T(n) = \sum_{i=0}^{n-1} \sum_{j=i+1}^n 1 = \sum_{i=0}^{n-1} (n - i) = \frac{n(n + 1)}{2} \approx \frac{n^2}{2}$$

# Example

If  $n$  is the number of items in the array to sort

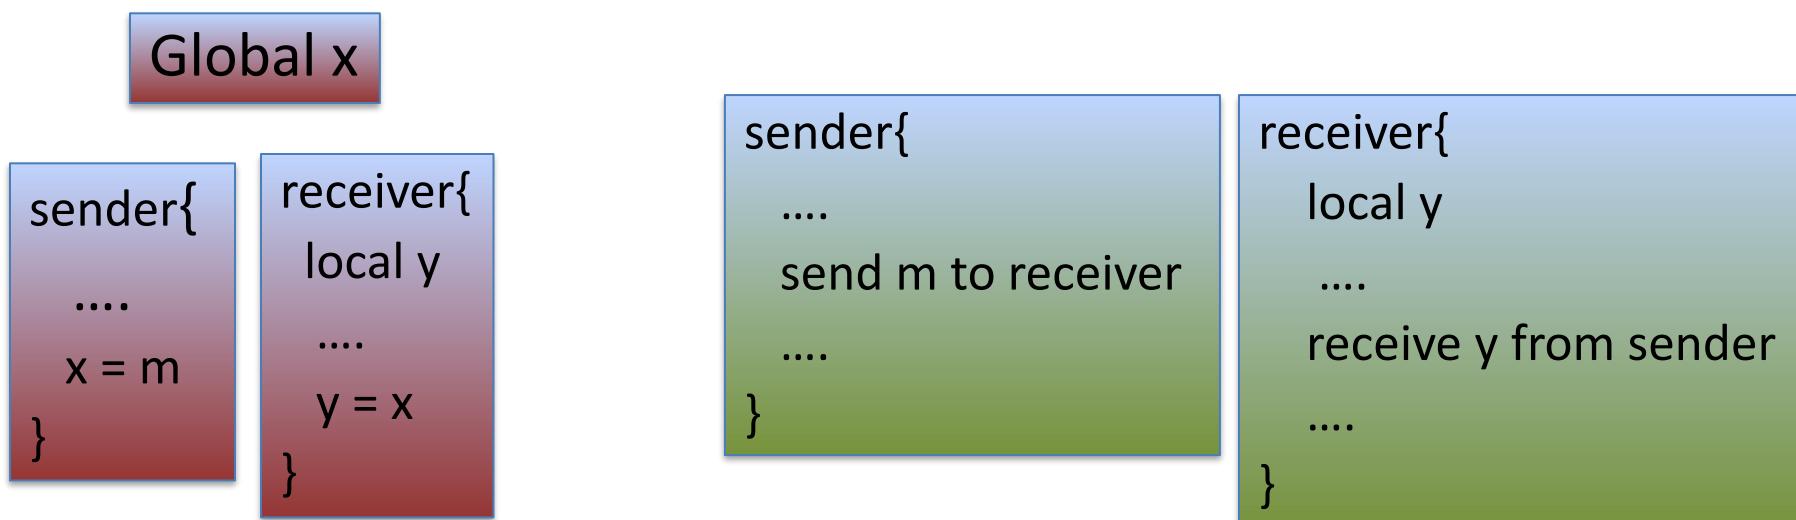
- 
- Case sequential:  $T(n) \sim n^2/2$
  - Case sequential with two invocations:  $T(n) \sim n^2/4 + n$
  - Case parallel with two invocations  $T(n) \sim n^2/8 + n$

$n \backslash T(n)$	$n^2/2$	$n^2/4+n$	$n^2/8+n$
20	200	120	60
40	800	440	220
1000	500000	256000	125500

# Communication and Synchronization

The processes involved in a concurrent program usually communicate among them and synchronize their actions

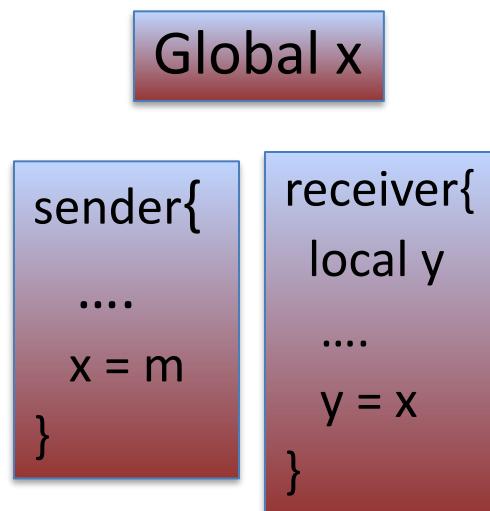
- ▶ Communication may be performed by means of shared memory
- ▶ ... or by some way of message passing



# Communication and Synchronization

The processes involved in a concurrent program usually communicate among them and synchronize their actions

- ▶ Communication may be performed by means of shared memory



This is an example of synchronization condition

The receiver **can not read** the variable (x) **until** the sender has written it.

# Problems with Concurrent Programming

## Atomic sentences

- ▶ Critical section
- ▶ Mutual exclusion
- ▶ Example: Tickets vending machine

Tickets can be bought at the same time in different ticket offices



# Problems with Concurrent Programming



- Each ticket office is a process  $T_i$
- All the ticket offices work concurrently
- All the ticket offices execute the same code
- The seats states of the theatre are stored into an array of booleans (taken or not)

```
seats= new Array[Boolean](200)
```

```
T1 || T2 || ... || Tn
```

## Code of T1

```
while (true){  
    Show untaken seats to the user  
    a = ... // selected seat  
    seats[a] = true  
    print the ticket  
}
```

.....

## Code of Tn

```
while (true){  
    Show untaken seats to the user  
    a = ... // selected seat  
    seats[a] = true  
    print the ticket  
}
```

# Example: Tickets Vending Machine

- ▶ Such a solution is incorrect. May happen **two customers buy tickets for the same seat**

Ti: show untaken seats to the user

Tj: show untaken seats to the user

Ti: `a = ... // selected seat in office i, for example a = 22`

Tj: `a = ... // selected seat in office j, for example a = 22`

Ti: `seats[22] = true`

Tj: `seats[22] = true`

Ti: print the ticket

Tj: print the ticket



**Interleaving:** This is an execution trace, i.e. a branch of the tree which behaves incorrectly. A trace like this is useful to show that a program is **incorrect**

## Code of each office

```
while (true){
```

    Show untaken seats to the user

`a = ... // selected seat`

`seats[a] = true`

    print the ticket

```
}
```

The seat nº 22 has been assigned to two different customers

# Example: Tickets Vending Machine

We try to debug the code

## Code of Ti

```
while (true){  
    success= false;  
  
    while (!success){  
        Show untaken seats to the user  
        a=...; // selected seat  
        if (! seats[a]) {  
            seats[a] = true ;  
            print the ticket  
            success= true;  
        } else {  
            display an error message  
        }  
    }  
}
```

We include this code to be  
sure of a seat is actually  
untaken when selected

# Example: Tickets Vending Machine

## Code of Ti

```
while (true){  
    success= false;  
  
    while (!success){  
        Show untaken seats to the user  
        a=...; // selected seat  
        if (! seats[a]) {  
            seats[a] = true ;  
            print the ticket  
            success= true;  
        } else {  
            display an error message  
        }  
    }  
}
```

### Execution TRACE which shows the error

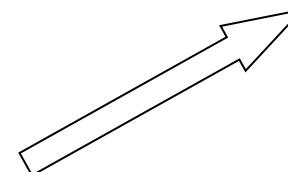
```
Ti: check seats[22] as untaken  
Tj: check seats[22] as untaken  
Ti: seats [22] = true  
Tj: seats [22] = true  
Ti: print the ticket  
Tj: print the ticket
```

We have the same problem as before,  
but now we have found out the lines of  
code where the problem is

# Example: Tickets Vending Machine

- **Atomic sentences:** Those that are executed with **no interruption**
- Only machine instructions are really atomic

```
while (true){  
    success= false;  
  
    while (!success){  
        Show untaken seats to the user  
        a=...; // selected seat  
        if (! seats[a]) {  
            seats[a] = true ;  
            print the ticket  
            success= true;  
        } else {  
            display an error message  
        }  
    }  
}
```



load seats[a] to CPU register  
test the value of CPU register  
jump to L1 if true  
set seats [a] to true  
code for "print the ticket"  
set success to true  
L1: code for "error message"  
L2: code following if statement

What would happen  
if the highlighted  
code were atomic?

>  
load seats[a] to CPU register  
test the value of CPU register  
jump to L1 if true  
set seats[a] to true  
>  
code for "print the ticket"  
set success to true  
L1: code for "error message"  
L2: code following if statement

# Example: Tickets Vending Machine

- Atomic sentences: Those executed with no interruption

**Only machine instructions are really atomic**

If the highlighted code were atomic, then an erroneous trace could not happen

Execution TRACE which shows the error

Ti: check seats[22] as untaken

**Tj: check seats[22] as untaken**

Ti: seats [22] = true

**Tj: seats [22] = true**

Ti: print the ticket

Tj: print the ticket

>  
load seats[a] to CPU register  
test the value of CPU register  
jump to L1 if true  
set seats [a] to true

>  
code for "print the ticket"  
set success to true  
L1: code for "error message"  
L2: code following if statement

# Example: Tickets Vending Machine

- Atomic sentences: Those executed with no interruption

**Only machine instructions are really atomic**

Valid interleavings are only those which happens **BEFORE OR AFTER** Ti has executed the highlighted code

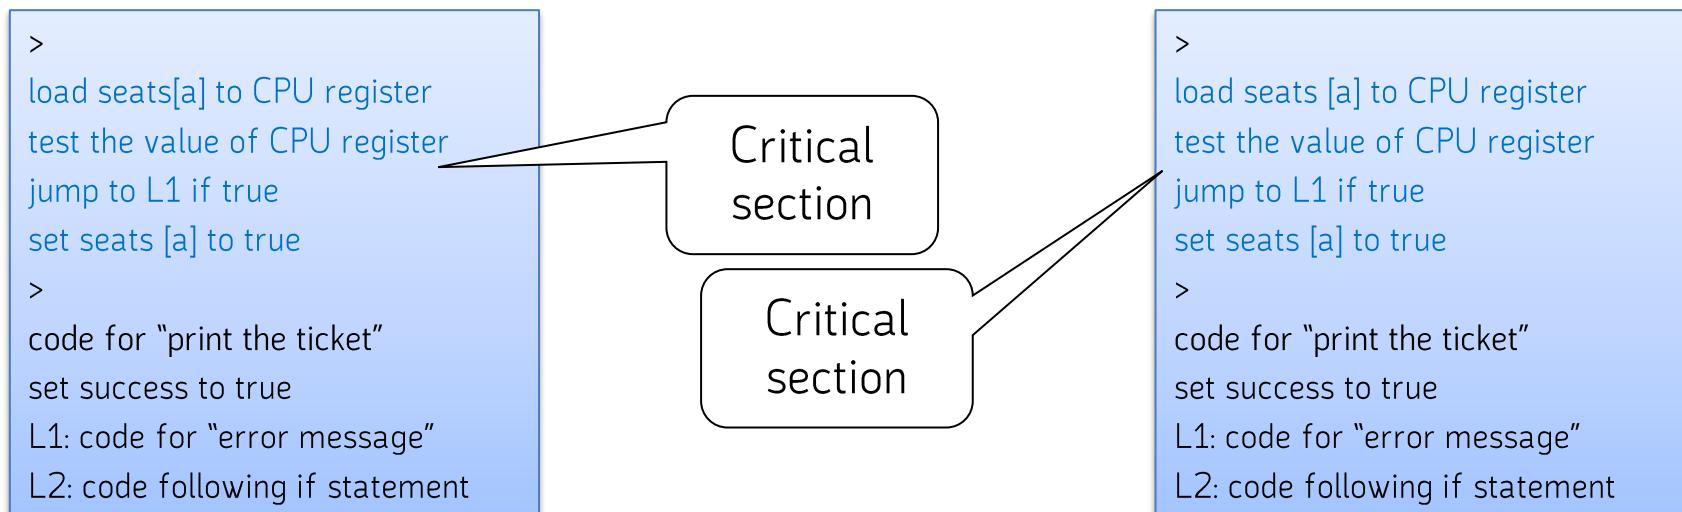
Valid execution TRACE

Ti: check seats[22] as untaken  
Ti: seats[22] = true  
Tj: check seats[22] as **taken**  
...

>  
load seats[a] to CPU register  
test the value of CPU register  
jump to L1 if true  
set seats [a] to true  
>  
code for "print the ticket"  
set success to true  
L1: code for "error message"  
L2: code following if statement

# Example: Tickets Vending Machine

- Atomic sentences: those executed with no interruption
  - Only machine instructions are really atomic
- **Critical section**: a block of code in a process that ought to be executed as atomic
- When two critical sections of two different processes can not be overlapped when executed (interleaving is not allowed) then it is stated that they must be executed with **mutual exclusion**



- If  $CS_i$  and  $CS_j$  are critical regions of the processes  $T_i$  and  $T_j$ , then a mutual exclusion states  $\forall e. (CS_i \rightarrow CS_j) \vee (CS_j \rightarrow CS_i)$

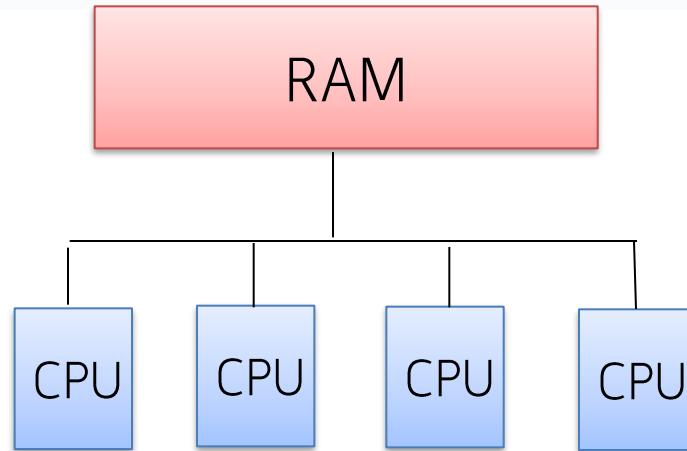
# Platforms

- **Single processor systems**

- Concurrency only can be implemented by interleaving the sentences of the processes
- It is useful to serve many users with a single computer
- When a process executes blocking I/O operations, the idle CPU cycles can be used to execute other processes
- The memory is shared by all processes, so communication is straightforward by means of shared memory
- Communications by means of message passing are also allowed

# Platforms

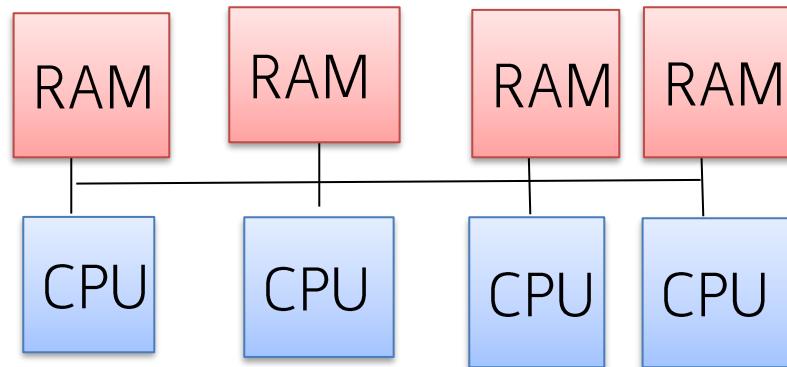
- Multiprocessors highly coupled



- A true parallelism can be achieved
- The natural communication among processes is performed by means of shared memory
- Each processor may have its own local memory (usually faster)
- Usually, each CPU is assigned more than one process
- This architecture has a high performance due to the parallel execution and the low cost of communication operations

# Platforms

- Multiprocessors loosely coupled (Distributed Systems)



- This architecture allows real concurrency
- Each node of this network may be a single processor or a multiprocessor highly coupled
- The natural communication among processes is performed by means of message passing
- In this architecture communications take much time. Even more, this may lead to neutralize the performance gained due to the parallel execution